

COMP 540 Recitation

Week 12: Finessing the Term Project

TAs: Gunny, Issac, Raymond, Tian

Tasks

- Recommended Timeline
 - Project, Homework 6, Final Exam
- Baseline Recommendations Walk-Thru
 - Pipeline
 - Pre-processing: Augmentation
 - Modelling
 - Ensembling
 - Report & Presentation
- Bonus Recommendations Walk-Thru
 - Bonus Models
 - Bonus Post-processing
- Workflow Recommendation

Timeline

- **ConvNet** -> less than 1 hour's Training Time
- **U-Net** -> less than 1 hour's Training Time
- **Masked Regional ConvNet** -> 24 hours Training Time (Purported)
- Pac-Man Reinforcement Learning (Homework 6) -> 1 Afternoon's Training Time
- Apr 18 -> Project Presentation
- Apr 20 -> Homework 6 Due Date
- Apr 30 -> Final Exam
- Recommendation:
- **Finish your preprocessing-modelling-ensembling work next week**

Example Pipeline


1. Set up k fold validation
2. Set up data augmentation
3. Train your model
4. Perform test time augmentation
5. Feed results into meta learner
6. Submit!

Baseline: Data Augmentation

- Not all training samples look the same
- You have to normalize them
- Numerically
- Stylistically
- Elastic Transform (affine transformation that involves “local distortion” of pixels e.g. making MNIST look like shaky handwriting)
- Averaging predictions (done during training and test time)

Baseline: Models

U-Net




Baseline: Ensembling

- Majority Vote
 - Extremely basic: examining majority vote
- Bagging
 - essentially taking the mean of individual model submissions
- Ensembling highly uncorrelated, but weaker performing models is better!
- <https://mlwave.com/kaggle-ensembling-guide/>

Baseline: Report & Presentation

- Baseline Expectations:
- Explanation of Justification of Pre-processing Techniques
- Explanation and justification of Algorithms
- Explanation and Justification of Ensembling Techniques
- Explanation and Justification of Basic Machine Learning Techniques
 - Hyperparameter Selection
 - Bias & Variance Tradeoff
 - Cross-Validation
 - Overfitting & Underfitting
 - Solution to Problem of Generalization

Object Detection



car

car

bus

car

car

car

traffic light

car


person

person

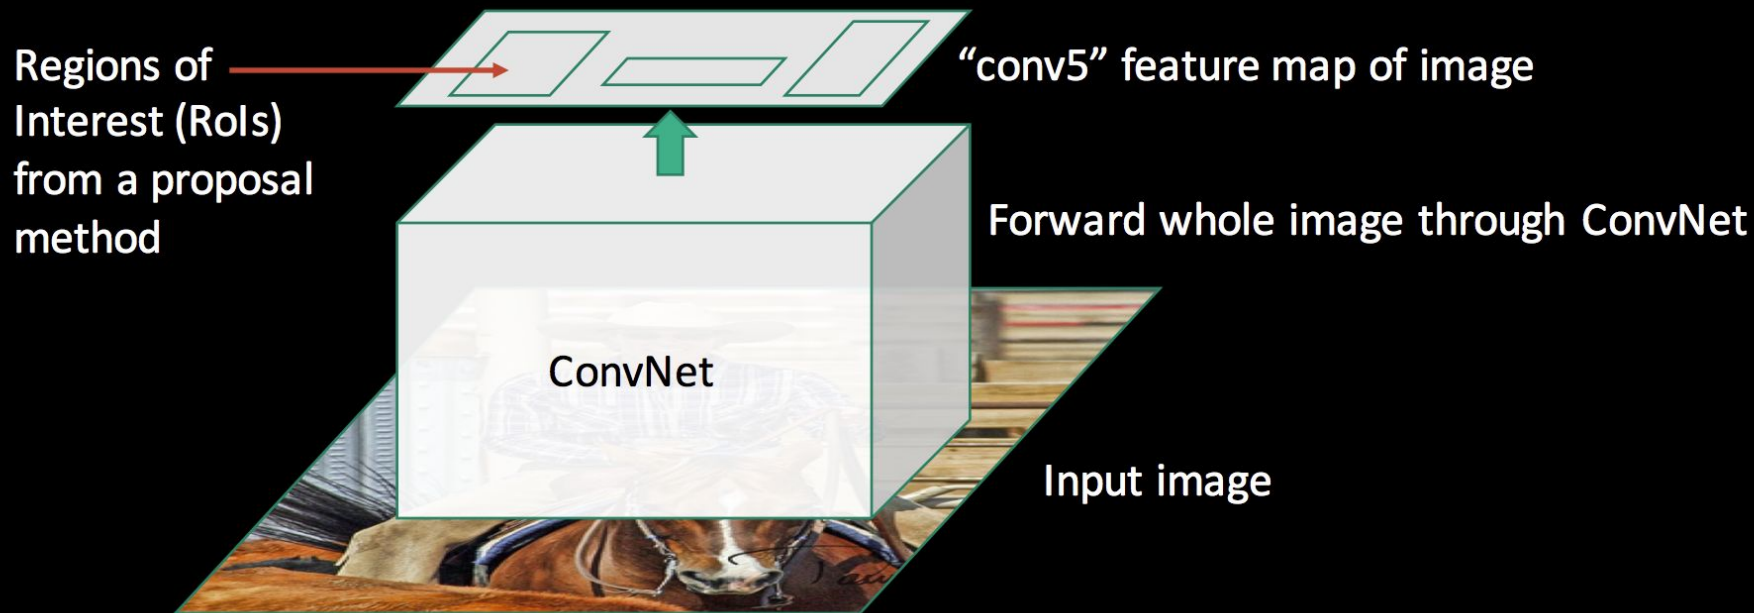
person

handbag


RCNN




Faster-RCNN




Faster-RCNN

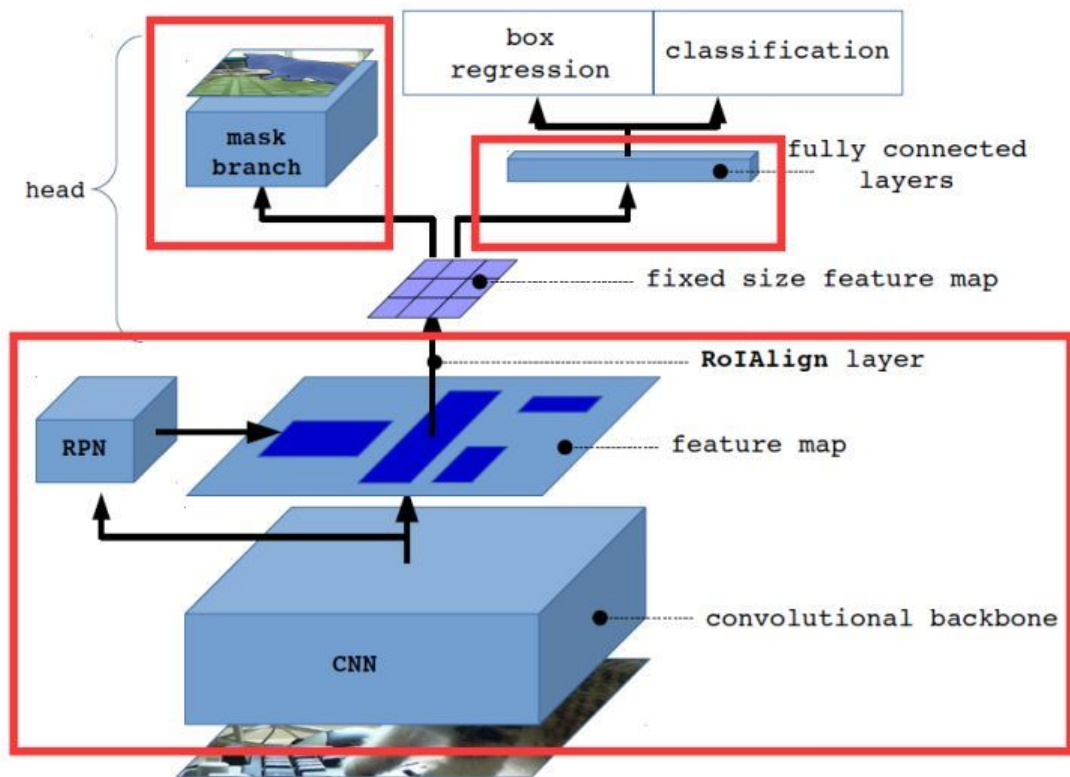


VS





Mask-RCNN



Other State-of-the-Art Models

- YOLO
- **Single-shot, Multi-object Detection**

Bonus: Super Learner

Set up the ensemble

- Specify a list of L base algorithms (with a specific set of model parameters).
- Specify a metalearning algorithm.

Train the ensemble

- Train each of the L base algorithms on the training set.
- Perform k -fold cross-validation on each of these learners and collect the cross-validated predicted values from each of the L algorithms.
- The N cross-validated predicted values from each of the L algorithms can be combined to form a new $N \times L$ matrix. This matrix, along with the original response vector, is called the "level-one" data. (N = number of rows in the training set)
- Train the metalearning algorithm on the level-one data.
- The "ensemble model" consists of the L base learning models and the metalearning model, which can then be used to generate predictions on a test set.

Bonus: Super Learner

Predict on new data

- To generate ensemble predictions, first generate predictions from the base learners.
- Feed those predictions into the metalearner to generate the ensemble prediction.

Bonus: Post-Processing

- Note that your test score is evaluated as **masks around individual cells**, not an entire mask on a input scan sampe
- This could lead to a double loss in your score
- Use classic computer vision techniques (for e.g. edge detection) to segment and identify masks blobs for individual cells

Recommendations: Workflow

- Amazon Web Service EC2 | p2.xlarge
- <https://aws.amazon.com/ec2/instance-types/p2/>
- Jet's CUDA-CuDNN-Keras-Torch Bootstrap Installer
- <https://github.com/abhmul/InstallCUDA-Kerasv2>
- Dev Kit: Keras, numpy, jupyter, skimage
- https://chrisalbon.com/software_engineering/cloud_computing/run_project_jupyter_on_amazon_ec2/
- Advanced Dev Kit: pytorch
- Masked R-CNN, U-Net: Use Working Implementations with Citations
- Technical Blogs for Bonus Players

Recommendations: Workflow

- **Stop instance**, don't terminate instance
- This saves disk state on your AWS virtual machine
- Avoid developing library code for U-Net and Masked R-CNN
- It's great if you want application development practice, but a vanilla Masked R-CNN implementation would require a week to write and test
- You are graded on:
 - Application & Explanation of Models
 - Application & Explanation of Pre/post-processing, Ensembling
 - Make sure you explain them in text (report) and verbally (presentation)
 - Avoid reading technical papers. **Blogs and kaggle discussion boards** are more time-efficient and human-friendly